

Yash Sharma

Email: yashsharma37158@gmail.com | Phone: +91-7974906887

LinkedIn: linkedin.com/in/yash-sharma-7b688a19b | GitHub: github.com/yash37158

SUMMARY

AI engineer with 2+ years building production systems for enterprise and startup clients in Singapore and India. Builds with LLMs and agents - MCP servers, Gemini workflows running in production - and applies AI-assisted development across live client delivery. Full-stack foundation in Go, TypeScript, REST APIs, Kubernetes and cloud.

SKILLS

AI & Agents: MCP servers, AI agents and orchestration, LLM integration (Gemini), vLLM inference observability, prompt engineering, AI-assisted development (Claude Code, Cursor, GitHub Copilot), AI-assisted code review, AI-driven automation (Zapier, n8n)

Languages: JavaScript/TypeScript, SQL, Python (basic)

Backend & Data: REST API design, Node.js, microservices, event-driven systems (Kafka), real-time integrations, PostgreSQL, MongoDB, ClickHouse

Frontend: React, Next.js, Vue.js

Cloud & DevOps: AWS, Huawei Cloud, Docker, Kubernetes, GitHub Actions, CI/CD

Quality & Delivery: Automated testing (Cypress, Playwright), CI pipeline automation, code review, Agile/Scrum

EXPERIENCE

Contract Full Stack Engineer | June 2025 - Present

EMedia-Plus Pvt. Ltd., Singapore | Dec 2025 - Present

- Modernized a legacy enterprise platform, shipping 90+ REST APIs and a template-based Vue.js CMS, by leading solution design for a Delphi-to-Go migration with AI-assisted development (Claude Code, Cursor) applied across authoring, refactoring and code review.
- Sustained client SLA adherence across hundreds of customer-managed Android terminals, by building an offline-capable player with resumable 5GB uploads, encrypted real-time content delivery, and containerized microservices on Kubernetes across AWS and Huawei Cloud.
- Reduced repetitive engineering and delivery effort, by building AI agents and automation workflows with MCP tooling, Zapier and n8n and wiring them into the team's day-to-day process.

TripPro, Delhi | Sept 2025 - Nov 2025

- Automated personalized content generation at 100+ outputs per day with sub-second latency, by designing and deploying Gemini-based LLM workflows end to end, from prompt design through cloud release.
- Reduced data-layer latency 30%, by analyzing MongoDB aggregation pipelines and redesigning indexing strategy and React interfaces.

Honestlee (Startup), Remote | June 2025 - Sept 2025

- Compressed feature delivery from weeks to days as sole engineer, by converting ambiguous stakeholder requirements into scoped tasks and shipping REST APIs, authentication and service modules in Node.js, PostgreSQL and MongoDB with AI-assisted development.
- Delivered customer onboarding, role-based access control and geolocation workflows for every new user, by building them in React with Mapbox and integrating Zoho CRM over its API.

Open Source Engineer (Meshmate) - Layer5, Meshery (CNCF Project) | Oct 2023 - June 2025

- Shortened CI feedback loops across a globally used open-source platform, by migrating the full test suite from Cypress to Playwright and automating pipelines with GitHub Actions and Docker.
- Raised new-contributor onboarding success across a globally distributed community, by writing documentation, running boot camps and mentoring contributors alongside maintainers.

Full Stack Developer (MERN) - Eskay's, Mumbai, India | Jul 2022 - Jan 2023

- Reduced page load latency 47%, by rebuilding the checkout experience on Stripe with targeted front-end and server-side performance optimization.
- Delivered consistent real-time order processing for every customer, by integrating third-party logistics and mapping services (Shiprocket, Google Maps) with Firebase and MongoDB Atlas.

PROJECTS

- **ScaleMind** | [Link](#) - MCP (Model Context Protocol) server that lets DevOps engineers manage cloud infrastructure scaling through natural language conversations.
- **Orchestr8 (GPU & Inference Observability)**, | [Link](#) - Multi-tenant platform in Go, ClickHouse and OpenTelemetry correlating NVIDIA DCGM hardware telemetry with vLLM inference latency to separate GPU thermal throttling from traffic surges, validated by automated fault injection on a live Kubernetes cluster and hardened by 44 adversarial security tests.
- **Meshify (Infra Dashboard)**, | [Link](#) - Kubernetes-native observability dashboard built in Go and React with Istio and Cilium, designed for infrastructure visibility and operational troubleshooting.

EDUCATION

Master of Computer Applications (MCA) - Amity University, Mumbai, India | 2023 - 2025

Bachelor of Computer Applications (BCA) - Amity University, Mumbai, India | 2020 - 2023